

276: INTRODUCTION TO PROLOG

Exercise 1: ‘Introduction to Prolog’

14 January 2016

Objectives

The objective of this laboratory session is to get you started with the programming language Prolog and the Sicstus Prolog environment.

I. Introduction to Sicstus Prolog

1. To start the Sicstus environment type the following command into a terminal window:

```
sicstus
```

The prompt `| ?-` indicates that the environment is ready to receive commands from you, such as loading a program or executing a query.

2. Enter the following query at the Sicstus prompt.

```
write('Hello world').
```

Make sure there is no space between `write` and the opening `(`. That is a syntax error. (Try it.) Spaces elsewhere are fine. Do not forget the `'.'` at the end. Sicstus will just start a new line and wait for more input. Every query and every clause in Prolog must end with `'.'`.

3. Sicstus should evaluate the query as soon as you press Enter. In this case, the built-in predicate (command) `write/1` is executed. This predicate is ‘extra-logical’. That is to say, it always succeeds: its purpose is to print out its argument. You will also see that Sicstus outputs a `yes` to confirm that the query succeeded. For comparison, enter the query

```
integer(a).
```

which should confirm that `a` is not an integer.

4. Exit Sicstus by executing the query `'halt.'` or by pressing `Ctrl-D`.
5. Simon Coffey (spc03) writes: After about eight seconds of using Sicstus under Linux, you will notice that it doesn’t do command-line history or editing. Four seconds later, this will make you want to break something. Don’t. Installed on the lab machines is a utility called `rlwrap` which will make your life much easier. If you run sicstus with the command:

```
rlwrap sicstus
```

command-line editing will be enabled. It even remembers the commands you used last session. If you add the following line to your `.cshrc` file (located in your home directory), you can launch it as normal by just typing `sicstus`

```
alias sicstus "rlwrap sicstus"
```

Loading programs

- This exercise uses a simple Prolog program `sample.pl`. Download a copy from CATE. It is normal to work with two windows simultaneously: one for the editor, the other for Sicstus. So: exit Sicstus by typing `'halt.'` or pressing `Ctrl-D`, and launch the editor (in the background) and Sicstus together by typing:

```
gedit sample.pl &
sicstus
```

(You can use any other text editor if you have a preferred one.)

Alternatively (better): open two terminal windows and use one for Sicstus and the other one for everything else (editors, examining directories and files, etc.).

You can use whichever text editor you like. Some support Prolog syntax highlighting.

For `emacs` users, there is a Prolog interface. See the Sicstus manual for details.

- Load the program into the Sicstus environment, by typing either of: `compile(sample).` or `[sample].` `'[sample].'` is just a shorthand version. This will look for a file named `sample.pl` in the current working directory. Remember that all the Prolog commands and clauses must finish with a dot (`.`).

You can also start Sicstus with the `-l` option on the command line. Type

```
sicstus -l sample.pl
```

- Try the following queries or commands:

```
check_number( 0 ).
check_number( 100 ).
check_number( 531 ).
```

(The spaces around the numbers are just for readability.)

What is wrong with the last query? It loops! To interrupt its execution type `Ctrl-C`. You'll get the prompt

```
Prolog interruption (h for help)?
```

Reply with `a`, for abort, in order to stop the execution.

- The reason for the infinite loop is the lack of a terminating case for odd numbers. To fix this, add the clause

```
check_number( 1 ).
```

to be the first or second clause in your copy of the `sample.pl` file. Note that this modification has to be done in the window that is running the editor, and that you have to save the file, before returning to the window running Sicstus and reloading the file by typing:

```
compile(sample). or [sample].
```

- To exit the Sicstus Prolog environment use the command `halt`. Don't forget that all Prolog commands and clauses finish with a dot.

Debugging (tracing)

Don't write monolithic code with bugs and then debug it. Use good programming practice.

Tip: do not ignore warning messages (if any) generated when you load/compile your program. The most common bug — and very hard to spot — is usually a simple mistyping of a predicate name or variable name or constant. The warning messages often help you to spot that.

- Sicstus has a built-in debugger (tracer) which allows you to step through and inspect each subquery that is generated as the program executes. You will often find it instructive to look at what your code is doing even if it produces the right answers.
- To invoke the debugger enter the query

```
trace.
```

If you now type a normal query (try e.g. `check_number(8).` with the `sample.pl` program) the interpreter will stop after every (non built-in) subgoal. Press **Enter** to move it on to the next one.

- During tracing, to *skip* over the execution of a subgoal without examining it, type **s** before pressing **Enter**.
- During tracing, to *zip* to the end of the current computation, type **z** before pressing **Enter**.
- When you have finished, back at the top level, use `'notrace.'` to turn the debugger off.
- Tracing everything in detail can be slow and tedious. Setting 'spypoints' speeds things up if you are only interested in examining calls to specific predicates. For example

```
spy(child_of/2).
```

adds a spy point for the predicate `child_of`. The `/2` is the arity (number of arguments) of the predicate. Spypoints can also be set and removed during tracing by entering **+** and **-** when a call to the appropriate predicate appears in the trace. You can *zip* to the next call with a spypoint attached by typing **z** during tracing.

- mjs writes: I find tracing to be of limited use for debugging except sometimes for a quick check. I hardly ever bother with spypoints. Often, a more effective method is simply to add temporary `write` statements to your program to print out the values of variables as it executes.
- `compile` command and `-l` option invoke the compiler. (Obviously.) You can also load a program using `consult`. *Consulting* a program loads it into memory to be interpreted. The code will run more slowly (you won't notice the difference for the programs you will be writing in this course) but more tracing information is recorded. To see the difference, try

```
consult(sample).
```

and then trace `check_number(8).` again.

Some notes

- The Department of Computing has installed Sicstus Prolog on all lab machines, running under Linux and Windows. The assessed lab exercises will be marked under Linux, so it is essential that you ensure your submissions run properly on the lab Linux installation. You might need to use the `fromdos` utility for converting source files if you develop your programs using the Windows implementation.

Sicstus itself is available for download under a DoC student licence agreement if you would like to install it on your own computers. See 'Software sources' under the CSG pages at `/vol/csg/download/software/sicstus`.

- You can invoke the complete on-line manual of Sicstus Prolog (PDF format) by using the shell command `sicstus-manual`. There is also a version on the web (see below). Have a brief look through it to get an idea of the range of built-in predicates and utility libraries provided. There is no need to study it in detail.

Don't print the manual. It is huge. See

`http://www.sics.se/sicstus/docs/latest4/html/sicstus.html`

or type `sicstus-manual` at the command line.

- There are several other Prolog implementations besides Sicstus. A popular and stable alternative, preferred by some users, is SWI-Prolog. It is free, and can be downloaded from `www.swi-prolog.org` if you would like to install it on your own computers. YAP Prolog, also free, is the fastest implementation generally available. You can download it at `www.yap.sourceforge.net`.

Please note: although there should be very few differences between Sicstus and these other implementations of Prolog for the programs you will developing on this course, all assessed work is marked (automatically) using Sicstus on Linux.

II. Prolog Database

The rest of this exercise is based on a database representing the fragment of a family tree illustrated in Figure 1.




Figure 1: Family Tree

The database is the set of atomic facts contained in the file `family.pl`. Download a copy from CATE. Load it in Sicstus with `'consult(family).'` or `'[family].'`. The file contains clauses about three predicates:

```

child_of( X, Y )      "X is a child of Y"
male( X )             "X is male"
female( X )           "X is female"

```

For instance, the first fact in the database is

```
child_of( emmeline, frank ).
```

and expresses "emmeline is a child of frank".

1. Enter the following queries, finding all the solutions:

```
child_of( peter, irene ).
child_of( peter, emlyn ).
child_of( X, george ).
child_of( george, Y ).
child_of( X, X ).
child_of( X, Y ).
```

2. Add further Prolog clauses which define the following predicates:

<code>mother_of(M, X)</code>	"M is the mother of X"
<code>grandparent_of(GP, X)</code>	"GP is a grandparent of X"
<code>daughter_of(D, X)</code>	"D is a daughter of X"
<code>uncle_of(Unc, X)</code>	"Unc is an uncle of X"
<code>niece_of(N, X)</code>	"N is a niece of X"
<code>great_grandfather_of(Gfx, X)</code>	"Gfx is a great-grandfather of X"
<code>ancestor_of(Anc, X)</code>	"Anc is an ancestor of X"

For the last one, compare the clauses for `superiorOf/2` in the lecture notes.

Test each one by posing suitable queries and finding all their solutions. As a guiding example, this is how we might define "X is a sister of Y":

```
sister_of( X, Y ) :-
    child_of( X, Z ),
    child_of( Y, Z ),
    X \= Y,
    female( X ).
```

(This definition is for illustration: it treats step-sisters as sisters.)

Note that in Prolog `\=` means 'do not unify'. You could also use `\==` (backslash, equals, equals) which means 'not identical'. Don't. Although you see `\==` in many books it is a *horrible* 'extra-logical' thing. This is not what it was intended for. Don't use it.

(If you don't understand the difference, compare what you get when executing Prolog queries `a \= b`, `a \== b`, `X \= a`, `X \== a`, `X \= Y`, `X \== Y`.)

And compare:

```
?- X == Y, X = Y.
?- X = Y, X == Y.
```

`==` and `=` are horrible things.)